

SQL as a Programming Language

Implementing Distance-Vector Routing
with Recursive CTEs in DuckDB

Mehdi Hoseyni

Seminar: SQL is a Programming Language - Summer Semester 2026

Outline

1. Motivation and Research Question
2. Distance-Vector Routing Protocol
3. The Bellman-Ford Equation
4. SQL Recursive CTEs: The Challenge
5. DuckDB USING KEY: The Solution
6. Standard vs Keyed CTE Comparison
7. Implementation: The Algorithmic Core
8. Convergence Results
9. Network Failure and Re-convergence
10. Live Demonstration
11. Benchmark Results
12. Key Findings and Conclusion

Presentation duration: 20 minutes + discussion

Demo: Interactive Streamlit application (localhost:8501)

Motivation and Research Question

Research Question

Can SQL serve as a **first-class programming language** for complex algorithmic simulation?

Approach: Implement the Distance-Vector Routing protocol entirely within SQL using DuckDB's recursive CTEs.

Goal: Show that declarative logic with state-aware recursion achieves global optimization through local updates.

Why It Matters

- **Traditional view:** SQL is only for data retrieval
- **Our claim:** Modern SQL extensions enable full simulation of distributed systems
- **Implication:** Database engines can function as protocol validation environments

Thesis: Modern SQL, through `USING KEY`, can express and execute distributed routing algorithms using declarative recursion.

Distance-Vector Routing Protocol

Core Idea - “Routing by Local Knowledge”:

- Each router maintains a *distance vector* of costs to all destinations
- Routers share tables **only with immediate neighbors**
- Each router updates its table if a cheaper path is discovered
- Process repeats until **convergence**

Algorithm Steps:

T=0: Each router knows only direct neighbor costs

T=1: Share tables; apply relaxation rule

T=2+: Repeat until convergence

T=k: Fixed point - all paths optimal

Key Properties

Distributed	No central controller
Iterative	Updates in rounds
Asynchronous	Independent timing
Convergent	$\leq V -1$ rounds


Real-World Protocols

- **RIP** (v1, v2) - max 15 hops
- **EIGRP** - Cisco hybrid protocol
- **BGP** - Internet backbone routing

The Bellman-Ford Equation

$$d(u, v) = \min_{w \in \text{Adj}(u)} \{ \text{cost}(u, w) + d(w, v) \}$$

Example Network Topology



Direct link A→C costs 23, but the indirect path A→B→C costs only 3+2 = 5.

Step-by-Step Convergence

Round	Discovery from Router A
T=0	B=3, C=23 (direct neighbors)
T=1	C=5 via B (3+2 < 23), D=28 via C
T=2	D=10 via B (3+2+5 < 28)
T=3	No changes - converged

SQL Recursive CTEs: The Challenge

The Problem

Standard SQL recursive CTEs are **append-only**:

- `UNION ALL` adds new rows each iteration
- Cannot update or delete existing rows
- All possible walks are enumerated
- Final `GROUP BY ... MIN(cost)` extracts answer

Result: Exponential row growth on dense graphs. SQLite exhausts memory beyond 12 nodes.

What DVR Needs

A physical router's forwarding table:

- **In-place updates:** overwrite when better route found
- **Bounded size:** only $O(V^2)$ entries at any time
- **Auto-termination:** stop when no improvement
- **Per-key pruning:** one entry per (src, dst)

Core mismatch: Standard recursion builds a growing set of paths. DVR requires a fixed-size table refined iteratively.

DuckDB USING KEY: The Solution

WITH RECURSIVE ... USING KEY (from_node, to_node) - DuckDB's extension that designates columns as a unique key. During recursion, new rows with matching keys **replace** old rows instead of appending.

In-Place Updates

Keying on (from_node, to_node) maintains one entry per router pair. Cheaper paths overwrite old entries.

Bounded Memory Usage

Working set stays at $O(V^2)$ regardless of graph density. No walk explosion.

Automatic Fixpoint Detection

Recursion halts when delta set is empty. No explicit iteration bound needed.

Memory Usage	Walk Explosion	Termination	Post-Hoc Agg.
$O(V^2)$	None	Automatic	Not needed

Standard CTE vs Keyed CTE

SQLITE: STANDARD CTE

```
WITH RECURSIVE bellman(s,t,cost,hop)
AS (
  SELECT from_node, to_node,
         cost, to_node
  FROM edges
  UNION ALL
  SELECT b.s, e.to_node,
         b.cost + e.cost, b.hop
  FROM bellman b
  JOIN edges e ON b.t = e.from_node
  WHERE b.s <> e.to_node
)
-- POST-HOC aggregation required
SELECT s, t, MIN(cost), hop
FROM bellman GROUP BY s, t
```

DUCKDB: KEYED CTE

```
WITH RECURSIVE routing(s,t,cost,hop)
  USING KEY (s, t) AS (
  SELECT from_node, to_node,
         cost, to_node
  FROM edges
  UNION
  SELECT r.s, e.to_node,
         MIN(r.cost + e.cost),
         arg_min(r.hop,
                 r.cost + e.cost)
  FROM routing r
  JOIN edges e ON r.t = e.from_node
  WHERE r.s <> e.to_node
  GROUP BY r.s, e.to_node
)
SELECT * FROM routing
```

Key difference: Standard CTE uses `UNION ALL` (accumulates all walks), requiring final `GROUP BY`. DuckDB uses `UNION + USING KEY` (prunes in-place), keeping only the best cost per (s, t).

The Algorithmic Core

```
WITH RECURSIVE routing_table(
  from_node, to_node,
  best_cost, next_hop)
  USING KEY (from_node, to_node) AS (
    -- T=0: Direct neighbor discovery
    SELECT from_node, to_node,
           cost, to_node AS next_hop
    FROM edges
  UNION
    -- Bellman-Ford relaxation
    SELECT r.from_node, e.to_node,
           MIN(r.best_cost + e.cost),
           arg_min(r.next_hop,
                  r.best_cost + e.cost)
    FROM routing_table AS r
    JOIN edges AS e
      ON r.to_node = e.from_node
    LEFT JOIN recurring.routing_table
      AS ex ON ex.from_node = r.from_node
             AND ex.to_node = e.to_node
    WHERE r.from_node <> e.to_node
           AND (ex.best_cost IS NULL
                OR r.best_cost + e.cost
                  < ex.best_cost)
    GROUP BY r.from_node, e.to_node
  )
SELECT * FROM routing_table
ORDER BY from_node, to_node;
```

Line-by-Line

USING KEY: Stateful row updates

Base case: Direct neighbor costs

JOIN edges: Neighbor propagation

recurring.*: Accumulated FIB

IS NULL OR <: Relaxation condition

MIN + arg_min: Best advertisement

This single query replaces hundreds of lines of imperative routing code.

Convergence Results: Full Network


From	To	Cost	Next Hop
A	B	3	B
A	C	5	B
A	D	10	B
B	A	3	A
B	C	2	C
B	D	7	C
C	A	5	B
C	B	2	B
C	D	5	D
D	A	10	C
D	B	7	C
D	C	5	C

Observations

- **A → C = 5 (via B):** Indirect A-B-C (3+2) beats direct (23)
- **A → D = 10 (via B):** Multi-hop A-B-C-D
- **Convergence:** 3 rounds for 4 nodes
- All router pairs converge to the expected shortest paths

Shortest path A → C: A → B → C (cost = 5, hops = 2)

Forwarding Information Base Visualisation



Iterative Convergence (2 rounds)

From	To	d*	Next Hop
A	B	3	B
A	C	5	B
A	D	10	D
B	A	3	A
B	C	2	C
B	D	7	D
C	A	5	B
C	B	2	B
C	D	5	D

Convergence achieved after 2 iteration(s) — distance metrics minimized.

Network Failure and Re-convergence

Simulating Link Failure

```
DELETE FROM edges
WHERE (from_node='B' AND to_node='C')
OR (from_node='C' AND to_node='B');
```

Re-running the **same query** discovers next-best paths automatically.

Route	Before	After	Δ
A $\rightarrow$ C	5	6	+1
B $\rightarrow$ C	2	9	+7
C $\rightarrow$ B	2	9	+6



INTERACTIVE DEMONSTRATION

Live Demo

app

Stage 1 Network Setup

Stage 2 Routing Simulation

Stage 3 Delete Nodes

Stage 4 Reconvergence

Benchmark Compare

Algorithm Theory

DVR Evaluation Suite

DuckDB Keyed Recursion Analysis

Global Settings

Backend Engine

DuckDB (USING KEY)

Distance-Vector Routing Protocol Simulator

A comparative evaluation platform for distributed routing convergence. Simulating Distance-Vector convergence utilizing DuckDB's native key-pruning recursive CTEs, SQLite standard recursive query execution, and algorithmic Python reference implementations.

Phase 1

Topology Design

Initialize network topologies using standard benchmarks or custom edge specifications.

Phase 2

Convergence Analysis

Evaluate distributed Bellman-Ford execution step-by-step through SQL query logs.

Phase 3

Fault Injection

Induce link failures and node crashes to evaluate protocol resilience.

Phase 4

Dynamic Re-convergence

Analyze forwarding information base modifications post-topology disruption.

Phase 5

Empirical Performance


Compare execution latencies, memory footprints, and scalability under varying graph sizes.

Interactive Streamlit Simulator


Streamlit - DuckDB - Pyvis - Plotly

Benchmark Results

Standard benchmark results — DuckDB vs SQLite, CTE



Log-log plots comparing DuckDB vs SQLite (Standard CTE) — Complexity comparison



Key Findings

- DuckDB: **5-15x speedup** over SQLite
- SQLite: **infeasible beyond 12 nodes**
- DuckDB: **polynomial scaling** $O(V \cdot E)$
- Python: validated reference baseline

Backend	Complexity	Max N
DuckDB USING KEY	$O(V \cdot E)$	50+
SQLite Standard	$O(V \cdot E \cdot P)$	11
Python Bellman-Ford	$O(V^2 \cdot E)$	50+
Python Dijkstra	$O(V(V+E)\log V)$	50+
Python Floyd-Warshall	$O(V^3)$	50+

Key Findings

1. SQL as an Algorithmic Language

SQL can fully implement distributed routing protocols. A single recursive query replaces hundreds of lines of imperative code.

2. State-Aware Recursion

Eliminates exponential path explosion. Working set bounded at $O(V^2)$. Achieves 5-15x speedup over SQLite.

3. Automatic Re-Convergence

Link failure via SQL DELETE; re-running the same query discovers next-best paths. No code changes needed.

Central result: Modern SQL is no longer “just for queries.” With state-aware recursion, it is a first-class language for simulating and analyzing complex distributed protocols.

Conclusion and Future Work

Conclusion

- DuckDB's `USING KEY` provides a **1:1 mapping** between SQL recursion and router behavior
- Built a complete **interactive evaluation suite** with 6 phases
- Keyed CTEs maintain **polynomial scaling**; standard CTEs degrade exponentially
- Declarative approach: **more concise** and **verifiably correct**

Future Research

- **Path-Vector (BGP)**: SQL array types for AS-path tracking
- **Link-State (OSPF)**: Dijkstra via recursive SQL
- **Dynamic Traffic**: Real-time metrics in edges table
- **Large-Scale**: Internet AS-graph convergence

References:

- [1] R. Bellman, "On a Routing Problem," *Q. Appl. Math.*, 1958.
- [2] DuckDB Documentation, "Recursive CTEs with `USING KEY`," 2026.
- [3] A. S. Tanenbaum, *Computer Networks*, 6th ed., Pearson, 2021.

END OF PRESENTATION

Thank You

Questions and Discussion

Mehdi Hoseyni

Database Systems Seminar - University of Tübingen - Summer 2026

Backup Slides

Additional material for discussion

Complexity Comparison (Backup)

Algorithm	Time	Space	Distributed	Neg. Weights
Bellman-Ford	$O(V^2 \cdot E)$	$O(V^2)$	Yes	Yes
Dijkstra (heap)	$O(V(V+E) \log V)$	$O(V^2)$	No	No
Floyd-Warshall	$O(V^3)$	$O(V^2)$	No	Yes
DuckDB USING KEY	$O(V \cdot E)$	$O(V^2)$	N/A	No
SQLite Standard CTE	$O(V \cdot E \cdot P)$	$O(V^2 \cdot P)$	N/A	No

P = walk count: In SQLite, P denotes distinct walks enumerated before the final GROUP BY. For dense graphs, P grows exponentially with $|V|$.

CTE Semantic Comparison (Backup)

Aspect	Standard CTE (SQLite)	Keyed CTE (DuckDB)
Recursion model	Append-only (UNION ALL)	In-place update (UNION + KEY)
Working set	Grows each iteration	Bounded at $O(V^2)$
Row pruning	None during recursion	Auto per-key minimum
Aggregation	Post-hoc GROUP BY	Built into recursion
Termination	Explicit iteration bound	Automatic (empty delta)
Memory	$O(V^2 \cdot P)$, P = walks	$O(V^2)$
Walk explosion	High (exponential)	None (pruned by design)